

①

Week 9 : Anomaly Detectⁿ

- We'll try to model $p(x)$ from the given data.
- If $p(x) \leq \epsilon$, ~~the~~ the given input x will be considered an anomaly & $p(x) > \epsilon$ will correspond to normal behaviour.

Gaussian (Normal) distributⁿ

Suppose $x \in \mathbb{R}$ & $X \sim \mathcal{N}(\mu, \sigma^2)$ then,

$$p(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right) \quad \begin{array}{l} \mu \rightarrow \text{mean} \\ \sigma \rightarrow \text{std dev.} \end{array}$$

For a dataset $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$,

$$\mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

$$\sigma^2 = \frac{1}{m-1} \sum_{i=1}^m (x^{(i)} - \mu)^2 \quad \text{or} \quad \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu)^2$$

Density estimatⁿ

Suppose $x \in \mathbb{R}^n$.

$$p(x) = p(x_1; \mu_1, \sigma_1^2) p(x_2; \mu_2, \sigma_2^2) \dots \{p(x_n; \mu_n, \sigma_n^2)\}$$

(Assuming features x_1, x_2, \dots, x_n are indep.)

$$= \prod_{j=1}^n p(x_j; \mu_j, \sigma_j^2)$$

Evaluating an Anomaly detection algorithm:

- (i) We assume that we have labelled data i.e.
 $y=0$ if normal, $y=1$ if anomalous.

Splitting into training set, CV set & Test set

Suppose there are 10000 good engines, } labelled.
20 flawed engines }

Training set: 6000 good engines

CV set: 2000 good engines, 10 flawed engines

Test set: 2000 good engines, 10 flawed engines.

Possible evaluation metrics:

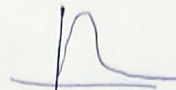
- (i) True positive, False pos, false neg, true neg
- (ii) Precision / Recall
- (iii) F1 score

* Classification accuracy isn't a good metric because of the skewed nature of the data.


Note: We can also choose ϵ using the CV set.

Anomaly detect ⁿ	Supervised learning
<ul style="list-style-type: none">- Small no. of +ve examples ($y=1$)- Large no. of -ve examples ($y=0$)- Hard for any algo to learn how the future anomalies look like given the small no. of +ve examples.	<ul style="list-style-type: none">- Large no. of both +ve & -ve examples.- Enough no. of +ve examples for an algo to get a sense of what +ve examples are like.
eg - Fraud detect ⁿ , Manufacturing, Monitoring machines in a data center.	eg - Spam classific ⁿ , weather predic ⁿ , cancer classific ⁿ .


Using the features to use

If feature x is a bit skewed,  then we can transform x to $\log x$ or $x^{1/2}$ or $x^{1/3}$ to get a more bell shaped distributⁿ.

$\log x$




$x^{1/2}$



Error analysis for anomaly detectⁿ

- Most common problem: $p(x)$ is comparable for both normal & anomalous examples.
- For ex: if there's a training example very close to a bunch of normal training examples but is actually abnormal, we should take a look at that ex. closely & come up with features specific to it

Multivariate Gaussian distⁿ



Sometimes when we model many gaussian distⁿs, the model fails to get the right boundary for the training set, so we use multivariate gaussian distⁿ.

$\mu \in \mathbb{R}^n$, $\Sigma \in \mathbb{R}^{n \times n}$ (Covariance matrix)

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp\left(-\frac{1}{2} (x-\mu)^T \Sigma^{-1} (x-\mu)\right)$$

$$\Sigma = \begin{bmatrix} a & c \\ c & b \end{bmatrix}$$

a, b - σ of x_1, x_2

c, d - covariance of x_1, x_2 if it's more then we get high prob. across $y=x$.

$$\mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

$$\Sigma = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu)(x^{(i)} - \mu)^T$$

- * Individual gaussian distributⁿs are special cases of multi-variate gaussian distributⁿs where the ellipses are axis aligned.

Original model	Multivariate model
<ul style="list-style-type: none"> - Have to manually create features to capture anomalies - Computationally cheaper - Should be used if m is small 	<ul style="list-style-type: none"> - Automatically captures anomalies in the features. - Computationally expensive - Should be used if $m \geq n$ or preferably $m \geq 10n$.

Recommender Systems

- Suppose there are m movies, u users who rate the movie.
 $r(i, j) = 1$ if user j has rated movie i .
 $y(i, j) = \text{rating value}$
- We want to find the ratings of movies as a particular user would have rated
- We basically train different linear regression algos for diff. users. & come up with $\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(u)}$.

Rating for a movie i : $(\theta^{(j)})^T x^{(i)}$ $x^{(i)} \rightarrow$ feature vector of movie i .
 by user j | $m(j) =$ no. of movies rated by user j .

~~min~~
min

$$\min_{\theta^{(j)}} \frac{1}{2m(j)} \sum_{i: r(i, j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i, j)})^2 + \frac{\lambda}{2m(j)} \sum_{k=1}^n (\theta_k^{(j)})^2$$

with change:

$$\min_{\theta^{(j)}} \frac{1}{2} \sum_{i: r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{k=1}^n (\theta_k^{(j)})^2$$

Over all $\theta^{(j)}$'s i.e. $\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(n_u)}$

$$\min_{\theta^{(1)}, \dots, \theta^{(n_u)}} \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i: r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2$$

Gradient Descent:

$$\theta_k^{(j)} := \theta_k^{(j)} - \alpha \sum_{i: r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) x_k^{(i)} \quad (k=0)$$

$$\theta_k^{(j)} := \theta_k^{(j)} - \alpha \left(\sum_{i: r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) x_k^{(i)} + \lambda \theta_k^{(j)} \right) \quad (k \neq 0)$$

Repeat this for all values of $j: 1, 2, \dots, n_u$.

Collaborative filtering

- Each user tells us their θ vector & we also know the ratings given by these users to some of the movies.
- We need to find an algo that comes up with the necessary features.

Algo:

Given $\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(n_u)}$, to learn $x^{(i)}$.

$$\min_{x^{(i)}} \frac{1}{2} \sum_{j: r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{k=1}^n (x_k^{(i)})^2$$

To learn $x^{(1)}, \dots, x^{(n_m)}$

$$\min_{x^{(1)}, \dots, x^{(n_m)}} \frac{1}{2} \sum_{i=1}^{n_m} \sum_{j: r(i,j)=1} (\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2$$

Procedure:

1. Guess θ
2. Learn features x
3. Get better estimate for θ .
4.

$$\theta \rightarrow x \rightarrow \theta \rightarrow x \rightarrow \dots$$

Collaborative filtering algorithm

$$\begin{aligned} J(x^{(1)}, x^{(2)}, \dots, x^{(n_m)}, \theta^{(1)}, \theta^{(2)}, \dots, \theta^{(n_u)}) \\ = \frac{1}{2} \sum_{(i,j): r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2 \\ + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2 \end{aligned}$$

- * Combining both optimising θ & optimising x .
- * minimise over $x^{(1)}, x^{(2)}, \dots, x^{(n_m)}, \theta^{(1)}, \theta^{(2)}, \dots, \theta^{(n_u)}$.
- * $x_0 = 1$ isn't used beoz the algo learns the features for itself.

Collaborative filtering algo:

1. Initialise $x^{(1)}, x^{(2)}, \dots, x^{(n_m)}, \theta^{(1)}, \theta^{(2)}, \dots, \theta^{(n_u)}$ to small random values.
2. Gradient descent for $j=1, 2, \dots, n_u, i=1, 2, \dots, n_m$.
$$x_k^{(i)} := x_k^{(i)} - \alpha \left(\sum_{j: r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) + \lambda x_k^{(i)} \right)$$
$$\theta_k^{(j)} := \theta_k^{(j)} - \alpha \left(\sum_{i: r(i,j)=1} (\theta^{(j)})^T x^{(i)} - y^{(i,j)} + \lambda \theta_k^{(j)} \right)$$
3. Predict: star rating = $(\theta^{(j)})^T x^{(i)}$.

Vectorizⁿ of algo: (Low rank matrix factorizⁿ)

$$Y = n_m \times n_u.$$

$$X = n_m \times n \quad \theta = n_u \times n$$

$$\text{Pred} = X \theta^T.$$

Finding related movies:

For each product i , we learn a feature vec $x^{(i)} \in \mathbb{R}^n$.
Movie j is similar to movie i , if

$$\|x^{(i)} - x^{(j)}\| \text{ is small.}$$

Mean normalizⁿ

- ~~Suppose~~ For each movie i , $x^{(i)} \rightarrow x^{(i)} - \mu^{(i)}$
Predict: $(\theta^{(j)})^T x^{(i)} + \mu^{(i)}$